

# Verification Condition Generation for eBPF

Troels Korreman Nielsen (xck773)

**Supervisor:** Ken Friis Larsen

PCC and eBPF  
Department of Computer Science  
University of Copenhagen

November 22, 2024

# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# What is eBPF?

eBPF allows Linux kernel extensions at runtime.

- Ensures high safety & performance.
- **Statically verifies** and **JIT-compiles** programs before running them.
- Used for network filtering, load balancing, monitoring, and security.

# eBPF Instruction Set

eBPF programs are submitted in the **eBPF ISA**:

- “Virtual machine”:
  - Eleven 64-bit registers.
  - Integer-addressed memory.
  - Ability to **call** external helpers.
- Heavily limited/restrictive:
  - Only 64-bit and 32-bit ALU operations.
  - No indirect jumps.
  - No routines\*.
  - No stack memory (only a frame buffer).
  - **Easy to validate**.

```
mov r1 52
add r1 3
div r1 3
mov r2 r1
mul r2 11
jlt r1 r2 label
mov r1 42
label:
mov r0 r1
exit
```

Example eBPF program.

All programs must be validated by the **verifier** before execution.

The verifier serves two goals:


- Goal #1: reject all unsafe programs.
- Goal #2: accept many safe programs.

Static analysis of programs is hard. We have a conflict between **ensuring soundness** and **improving completeness**.

## Safe programs must:

- Run to completion.
- Access memory correctly.
- Not leak secrets.
- Call helpers correctly.
- etc ...

# Complexity



- Accepting unsafe programs means **running unsanitized code in kernel space**.
- eBPF is a popular attack surface for privilege escalation.
- Many kernel exploits have stemmed from eBPF verifier unsoundness.<sup>1</sup>

---

<sup>1</sup>Several dozens of eBPF verifier errors tracked on [cve.mitre.org](https://cve.mitre.org)



# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# Proof-carrying code

- **Alternative approach** to program validation.
- The **producer** (user) supplies a formal **safety proof** along with each program.
- The **consumer** (Linux kernel) only needs to **check** the supplied proof.

PCC pushes **non-trivial** work to **userspace**.

# Proofs about programs

Proof system:

- Create a logic that will only derive safe programs.
- We model statement effects through **Hoare triples** of the form  $\{P\} S \{Q\}$ .  
Meaning: If precondition  $P$  is true, then postcondition  $Q$  is true after executing  $S$ .

Verification condition generation:

- From program, generate conditions that imply safety if proven.
- Separates the proof from the program.
- We use **weakest precondition calculus** to compute preconditions from postconditions:

$$WP : \text{Program} \times \text{Formula} \rightarrow \text{Formula} \quad \text{where} \quad \vdash \{WP(S, Q)\} S \{Q\}$$

# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# Verification condition generation (for eBPF)

Specified Hoare logic and VC-gen for:

- A **subset** of eBPF.
  - 64-bit arithmetic.
  - Conditional jumps.
  - Rudimentary memory.
- A **subset** of safety policies:
  - No division/modulo by zero.
  - Safe memory access.

Proof-of-concept implementation, written in **Rust**:

- Exports VCs to **WhyML** for solving.
- No proof extraction, no proof checking.

# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

- **Straightline** instructions can be handled using a WP-calculus:

$$WP(\text{assert } P, Q) = P \ \&\& \ Q$$

$$WP(\text{div } r \ v, Q) = v \neq 0 \wedge \forall v. v = \text{alu}(r, v) \implies Q[r \leftarrow v]$$

$$WP(\text{store } k \ m \ v, Q) = Q \wedge \exists p, s. \text{is\_buffer}(p, s) \wedge p \leq m < p + s - (k - 1)$$

...

- **Straightline** instructions can be handled using a WP-calculus:

$$WP(\text{assert } P, Q) = P \ \&\& \ Q$$

$$WP(\text{div } r \ v, Q) = v \neq 0 \wedge \forall v. v = \text{alu}(r, v) \implies Q[r \leftarrow v]$$

$$WP(\text{store } k \ m \ v, Q) = Q \wedge \exists p, s. \mathbf{is\_buffer}(p, s) \wedge p \leq m < p + s - (k - 1)$$

...

- But **control-flow** instructions cannot:

$$WP(\text{jlt } r0 \ r1 \ \text{label}, ???) = ???$$

We need a different approach for these.



# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)**
  - WP-calculus
  - Control graphs**
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# Control graphs

- All eBPF jumps and branches are statically specified (direct).
  - No routine calls.
- ⇒ Modules can be converted to a **control-flow graph** representation.

## Bubble sort

```
    mov r4 0
    mov r8 r2
    sub r8 7
L0:
    mov r3 8
L1:
    mov r5 r3
    add r5 r1
    ldxdw r6 [r5 - 8]
    ldxdw r7 [r5]
    jlt r6 r7 L2
    stxdw [r5 - 8] r7
    stxdw [r5] r6
L2:
    add r3 8
    jlt r3 r8 L1
    add r4 8
    jlt r4 r2 L0
    mov r0 0
    exit
```

# Control graph generation (split)

Split into blocks, separating:

- Before labels
- After jumps

```
mov r4 0
mov r8 r2
sub r8 7
```

```
L0:
mov r3 8
```

```
L1:
mov r5 r3
add r5 r1
ldxdw r6 [r5 - 8]
ldxdw r7 [r5]
jlt r6 r7 L2
```

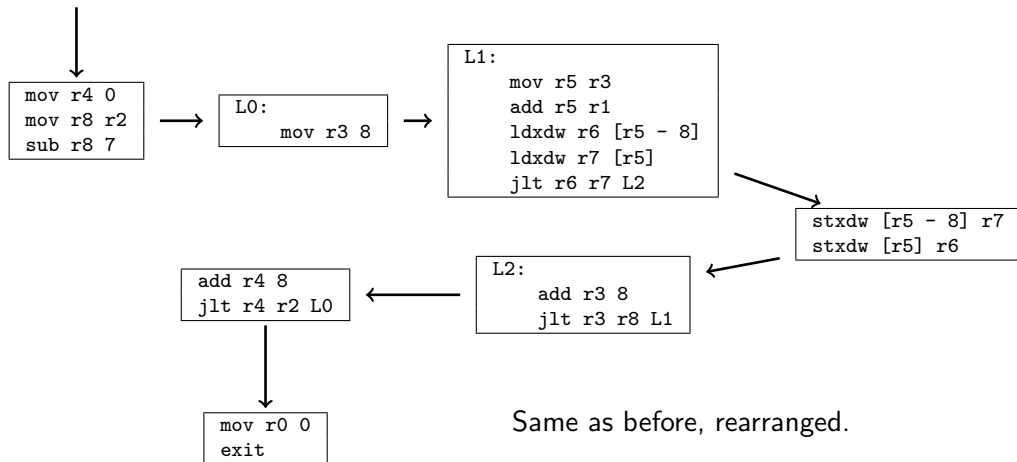
```
stxdw [r5 - 8] r7
stxdw [r5] r6
```

```
L2:
add r3 8
jlt r3 r8 L1
```

```
add r4 8
jlt r4 r2 L0
```

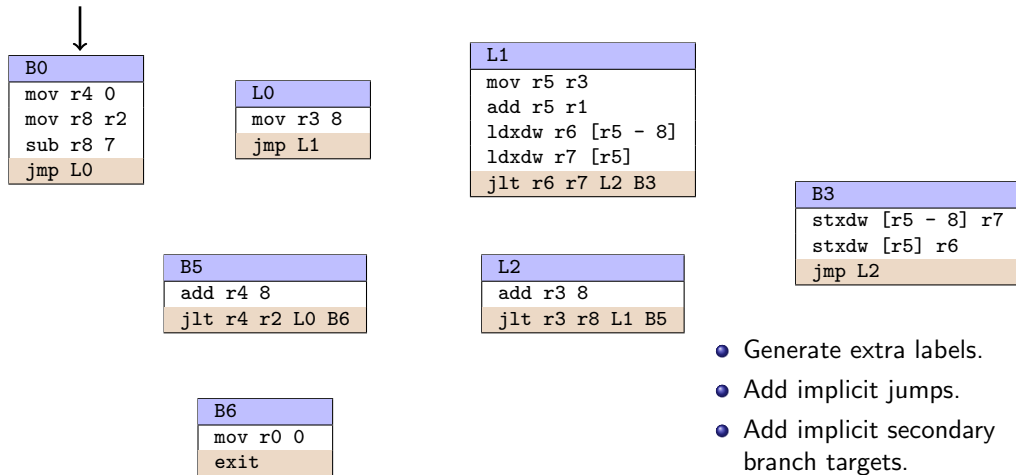
```
mov r0 0
exit
```

## Control graph generation (rearrange)



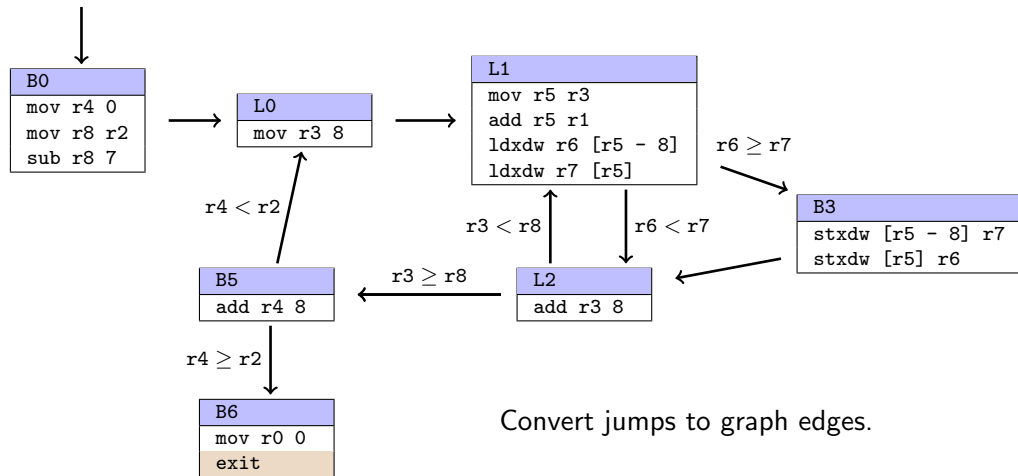
Same as before, rearranged.

# Control graph generation (blocks)



- Generate extra labels.
- Add implicit jumps.
- Add implicit secondary branch targets.


# Control graph generation (edge view)



Convert jumps to graph edges.

# Control graph

- View the control graph as a **dependency graph**.
- The **postcondition** for a block is created from the **conjunction** of its jump targets.




# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion




# Dependency resolution (cycles)

- In **cyclic** control graphs, we must contend with **dependency cycles**.
- A block might indirectly depend on its **own postcondition**.
- Prevents a finite definition of VCs.




# Dependency resolution (cycles)

Break up cycles by requiring  
user-annotated invariants.




# Dependency resolution (cycles)

Break up cycles by requiring  
user-annotated invariants.




# Dependency resolution (cycles)

Break up cycles by requiring  
user-annotated invariants.



# Dependency resolution (cycles)

Break up cycles by requiring  
user-annotated invariants.




# Dependency resolution (cycles)

A user-annotated block requirement:

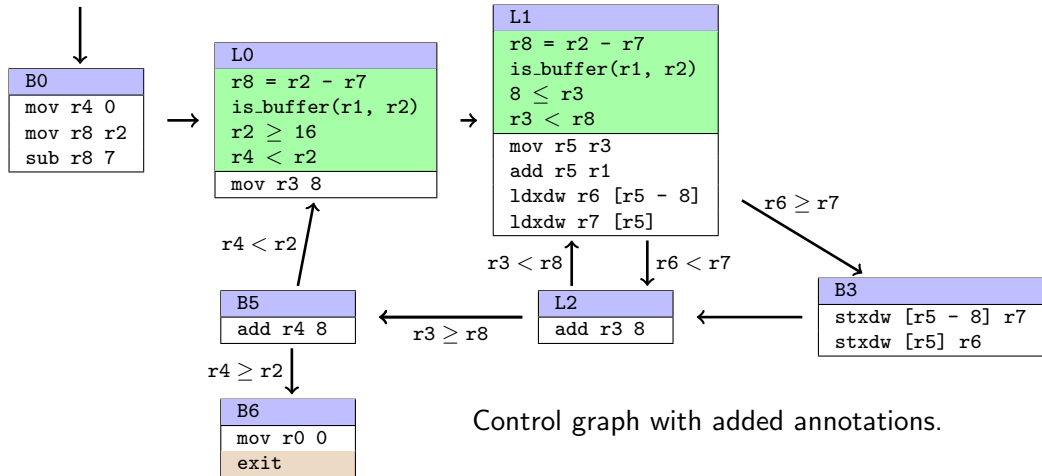
- Must be **strong enough** to validate its block.
- Must be **weak enough** to be validated by all dependants.

Annotations referred to by mapping:

$\alpha : \text{Label} \rightarrow \text{Formula}$



# Dependency resolution (cycles)



Control graph with added annotations.


# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)**
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)**
  - In practice
- 4 Termination
- 5 Conclusion



# Dependency resolution (traversal)


- Compute and cache in  $\Gamma$  the precondition for each block.
- Traverse blocks by **reverse topological order**.



# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$




# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$


$G_1 : \Gamma(B0)$



# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$




$G_1 : \Gamma(B0)$

$G_2 : \alpha_0 \implies \Gamma(L0)$

# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$



$G_1 : \Gamma(B0)$


$G_2 : \alpha_0 \implies \Gamma(L0)$

# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$


$G_1 : \Gamma(B0)$   
 $G_2 : \alpha_0 \implies \Gamma(L0)$



# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$
B6	$WP(B6, \text{true})$




$G_1 : \Gamma(B0)$

$G_2 : \alpha_0 \implies \Gamma(L0)$

# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$
B6	$WP(B6, \text{true})$
B5	$WP(B5, \alpha_0 \wedge \Gamma(L0))$



$G_1 : \Gamma(B0)$


$G_2 : \alpha_0 \implies \Gamma(L0)$



# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$
B6	$WP(B6, \text{true})$
B5	$WP(B5, \alpha_0 \wedge \Gamma(L0))$
L2	$WP(L2, \alpha_1 \wedge \Gamma(B5))$




$G_1 : \Gamma(B0)$

$G_2 : \alpha_0 \implies \Gamma(L0)$

# Dependency resolution (traversal)

(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$
B6	$WP(B6, \text{true})$
B5	$WP(B5, \alpha_0 \wedge \Gamma(L0))$
L2	$WP(L2, \alpha_1 \wedge \Gamma(B5))$
B3	$WP(B3, \Gamma(L2))$



$G_1 : \Gamma(B0)$

$G_2 : \alpha_0 \implies \Gamma(L0)$

# Dependency resolution (traversal)


(Ignore edge conditions for simplicity.)

$L$	$\Gamma(L)$
B0	$WP(B0, \alpha_0)$
L0	$WP(L0, \alpha_1)$
B6	$WP(B6, \text{true})$
B5	$WP(B5, \alpha_0 \wedge \Gamma(L0))$
L2	$WP(L2, \alpha_1 \wedge \Gamma(B5))$
B3	$WP(B3, \Gamma(L2))$
L1	$WP(L1, \Gamma(L2) \wedge \Gamma(B3))$

$G_1 : \Gamma(B0)$

$G_2 : \alpha_0 \implies \Gamma(L0)$

$G_3 : \alpha_1 \implies \Gamma(L1)$



# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

## In practice (input)

```

;# requires is_buffer(r1, r2)
;# requires r2 >= 16
    mov r4 0
    mov r8 r2
    sub r8 7
outer_loop:
;# req r8 = sub(r2, 7)
;# req is_buffer(r1, r2)
;# req r2 >= 16
;# req r4 < r2
    mov r3 8
inner_loop:
;# req r8 = sub(r2, 7)
;# req is_buffer(r1, r2)
;# req 8 <= r3
;# req r3 < r8
    mov r5 r3
    add r5 r1
    ldxdw r6 [r5 - 8]
    ldxdw r7 [r5]
    jlt r6 r7 skipped
    stxdw [r5 - 8] r7
    stxdw [r5] r6
skipped:
    add r3 8
    jlt r3 r8 inner_loop
    add r4 8
    jlt r4 r2 outer_loop
    mov r0 0
    exit
```

# In practice (output)

```
use mach.int.UInt64
use int.Int
use int.ComputerDivision
predicate is_buffer (p: uint64) (s: uint64)

goal inner_loop: forall r0 r1 r2 r3 r4 r5 r6 r7 r8 r9 r10 : uint64 . (((((r8 = (r2 - 7)) && is_buffer r1 r2) &&
(8 <= r3)) && (r3 < r8)) -> (forall v6 : uint64 . ((v6 = r3) -> (forall v5 : uint64 . ((v5 = (v6 + r1)) ->
((exists p3 : uint64 . (exists s3 : uint64 . (is_buffer p3 s3 /\ ((p3 <= (v5 + -8)) /\ ((v5 + -8) < ((p3 +
s3) - 7)))))) /\ (forall v4 : uint64 . ((exists p2 : uint64 . (exists s2 : uint64 . (is_buffer p2 s2 /\ ((p2
<= (v5 + 0)) /\ ((v5 + 0) < ((p2 + s2) - 7)))))) /\ (forall v3 : uint64 . (((v4 < v3) && (forall v2 :
uint64 . ((v2 = (r3 + 8)) -> (((v2 < r8) && (((r8 = (r2 - 7)) && is_buffer r1 r2) && (8 <= v2)) && (v2 < r8
))) \/\ (not ((v2 < r8)) && (forall v1 : uint64 . ((v1 = (r4 + 8)) -> (((v1 < r2) && (((r8 = (r2 - 7)) &&
is_buffer r1 r2) && (r2 >= 16)) && (v1 < r2))) \/\ (not ((v1 < r2)) && true)))))))))) \/\ (not ((v4 < v3)) &&
((exists p1 : uint64 . (exists s1 : uint64 . (is_buffer p1 s1 /\ ((p1 <= (v5 + -8)) /\ ((v5 + -8) < ((p1 +
s1) - 7)))))) /\ ((exists p0 : uint64 . (exists s0 : uint64 . (is_buffer p0 s0 /\ ((p0 <= (v5 + 0)) /\ ((v5
+ 0) < ((p0 + s0) - 7)))))) /\ (forall v2 : uint64 . ((v2 = (r3 + 8)) -> (((v2 < r8) && (((r8 = (r2 - 7))
&& is_buffer r1 r2) && (8 <= v2)) && (v2 < r8))) \/\ (not ((v2 < r8)) && (forall v1 : uint64 . ((v1 = (r4 +
8)) -> (((v1 < r2) && (((r8 = (r2 - 7)) && is_buffer r1 r2) && (r2 >= 16)) && (v1 < r2))) \/\ (not ((v1 < r2
)) && true))))))))))))))))))

goal outer_loop: forall r0 r1 r2 r3 r4 r5 r6 r7 r8 r9 r10 : uint64 . (((((r8 = (r2 - 7)) && is_buffer r1 r2) && (
r2 >= 16)) && (r4 < r2)) -> (forall v7 : uint64 . ((v7 = 8) -> (((r8 = (r2 - 7)) && is_buffer r1 r2) && (8
<= v7)) && (v7 < r8))))))

goal entry: forall r0 r1 r2 r3 r4 r5 r6 r7 r8 r9 r10 : uint64 . (((true /\ is_buffer r1 r2) /\ (r2 >= 16)) -> (
forall v10 : uint64 . ((v10 = 0) -> (forall v9 : uint64 . ((v9 = r2) -> (forall v8 : uint64 . ((v8 = (v9 -
7)) -> (((v8 = (r2 - 7)) && is_buffer r1 r2) && (r2 >= 16)) && (v10 < r2))))))))))
```

# Final step

In a proper PCC architecture:

- ① Pass to solver.
- ② Extract proof if successful solve.
- ③ Pass proof along with program to Linux subsystem.

In these experiments:

- ① Pass to Why3.
- ② Why3 tells us whether conditions can be proven or not.

✓ Why3 accepts the bubble sort.

# Comparison to verifier

VC-gen can validate things the verifier can't:

- Relationships between registers (and memory cells, theoretically).
- Non-contiguous value ranges.
- Dynamically sized buffers.
- ... and more.



# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination**
- 5 Conclusion

# Termination challenges

- If PCC for eBPF is to be viable, we need **total correctness**.

# Termination challenges

- If PCC for eBPF is to be viable, we need **total correctness**.
- Proof rules for termination in structured languages:

$$\frac{\{C \wedge P \wedge v = X\} s \{P \wedge v \prec X\}}{\{P\} \textbf{while } C \textbf{ invariant } P \textbf{ variant } v \textbf{ do } s \textbf{ done } \{P \wedge \neg C\}}$$

Where  $(\prec)$  is a well-founded relation.

# Termination challenges

- If PCC for eBPF is to be viable, we need **total correctness**.
- Proof rules for termination in structured languages:

$$\frac{\{C \wedge P \wedge v = X\} s \{P \wedge v \prec X\}}{\{P\} \textbf{while } C \textbf{ invariant } P \textbf{ variant } v \textbf{ do } s \textbf{ done } \{P \wedge \neg C\}}$$

Where ( $\prec$ ) is a well-founded relation.

- eBPF is **unstructured**, so above method doesn't apply.

# Termination challenges

- If PCC for eBPF is to be viable, we need **total correctness**.
- Proof rules for termination in structured languages:

$$\frac{\{C \wedge P \wedge v = X\} s \{P \wedge v \prec X\}}{\{P\} \textbf{while } C \textbf{ invariant } P \textbf{ variant } v \textbf{ do } s \textbf{ done } \{P \wedge \neg C\}}$$

Where  $(\prec)$  is a well-founded relation.

- eBPF is **unstructured**, so above method doesn't apply.
- Idea: Limit scope to “semi-structured” control graphs!

# Semi-structured control graphs

- “Semi-structured” control graph: partial ordering  $(\vec{<})$  of nodes in graph  $G$ .

# Semi-structured control graphs

- “Semi-structured” control graph: partial ordering  $\left(\overset{\rightarrow}{\lessdot}\right)$  of nodes in graph  $G$ .
- Straight path: Acyclic path from entry to node.

# Semi-structured control graphs

- “Semi-structured” control graph: partial ordering  $(\overset{\rightarrow}{<})$  of nodes in graph  $G$ .
- Straight path: Acyclic path from entry to node.
- Relations:
  - $A \overset{\rightarrow}{<} B$  if all straight paths to  $B$  include  $A$ .
  - $A \not\overset{\rightarrow}{<} B$  if no straight paths to  $B$  include  $A$ .







# Semi-structured control graphs

- “Semi-structured” control graph: partial ordering  $(\overset{\rightarrow}{<})$  of nodes in graph  $G$ .
- Straight path: Acyclic path from entry to node.
- Relations:
  - $A \overset{\rightarrow}{<} B$  if all straight paths to  $B$  include  $A$ .
  - $A \not\overset{\rightarrow}{<} B$  if no straight paths to  $B$  include  $A$ .
- Valid graph: for all edges  $A \rightarrow B \in G$ , either  $B \overset{\rightarrow}{<} A$  or  $B \not\overset{\rightarrow}{<} A$ .

# Semi-structured control graphs

- “Semi-structured” control graph: partial ordering  $(\overset{\rightarrow}{<})$  of nodes in graph  $G$ .
- Straight path: Acyclic path from entry to node.
- Relations:
  - $A \overset{\rightarrow}{<} B$  if all straight paths to  $B$  include  $A$ .
  - $A \not\overset{\rightarrow}{<} B$  if no straight paths to  $B$  include  $A$ .
- Valid graph: for all edges  $A \rightarrow B \in G$ , either  $B \overset{\rightarrow}{<} A$  or  $B \not\overset{\rightarrow}{<} A$ .
- Ordering is chronological:  
Edges jump either **forwards** or **backwards**.

# Examples



# Termination proofs

Any backwards jump targets a guaranteed previously visited node. We must then show:

- ① Decreasing variants for all backwards jumps.
- ② No sequence of backwards jumps can indefinitely increase a variant.

Not within scope of this presentation.

Nor fully within the scope of this project, for that matter.

# Table of Contents

- 1 eBPF
- 2 Proof-carrying code (PCC)
- 3 Verification Condition Generation (for eBPF)
  - WP-calculus
  - Control graphs
  - Dependency resolution (cycles)
  - Dependency resolution (traversal)
  - In practice
- 4 Termination
- 5 Conclusion

# Conclusion

In this project, I have:

- Identified safety policies for eBPF.
- Demonstrated VC-gen for an eBPF subset.
- Shown viability of VC-gen for branching & looping assembly.
- Shown PCC validating programs that Linux verifier cannot.

Future work:

- Proper exploration of termination proofs.
- Expansion of eBPF subset.
- Secret leak prevention.
- Proof-of-concept full framework.